


TRAINING: physics loss → backward through differentiable solver → same shared NCA weights



INFERENCE: one unseen layout → fixed physical representation → 64 local updates → topology